

Project Final Report

Report By:

Ramya Thadikonda -700762981

Aashik Shaik-700758163

Vaishnavi Beesa-700772902

Title:

Emotion-Based Movie Recommendation System Using Machine Learning and NLP

Abstract:

This project presents an Emotion-Based Movie Recommendation System that uses Machine Learning and Natural Language Processing to suggest movies according to a user's emotional state. Traditional recommendation platforms rely on ratings, user behavior, and trends, which do not reflect how a user feels at the moment. Our system bridges this gap by detecting user emotion from text input and classifying movies based on emotional tone using TF-IDF, Logistic Regression, and Naive Bayes. The recommendation engine then filters and ranks movies according to emotional relevance, popularity, and genre associations. The outcome is a personalized movie recommendation experience that aligns closely with user mood. Experimental results demonstrate classification accuracy between **85–90%**, proving the effectiveness of lightweight ML models for emotion detection.

1. Introduction

Background and Motivation

Recommender systems play a major role in online platforms such as Netflix, Amazon Prime, and YouTube. However, these systems mainly depend on user history, which does not reflect real-time user needs. Emotions greatly influence movie choices, making emotion-aware recommendations a logical next step.

Importance and Relevance

Users often spend a long time searching for movies that match their mood. An emotion-aware recommendation system helps reduce search time and improve user satisfaction. With increasing interest in AI-based personalization, integrating sentiment and emotion detection into recommendations is highly relevant.

Objectives and Research Questions

- Can we detect user emotion using NLP from text input?
- Can we classify movies into emotional categories using ML models?
- How effectively can we recommend movies based on emotional alignment?
- Can combining ML predictions with genre mapping improve personalization?

2. Problem Statement

Define the Problem

Traditional movie recommendation systems fail to consider the user's current emotional state. This results in irrelevant suggestions and poor user experience.

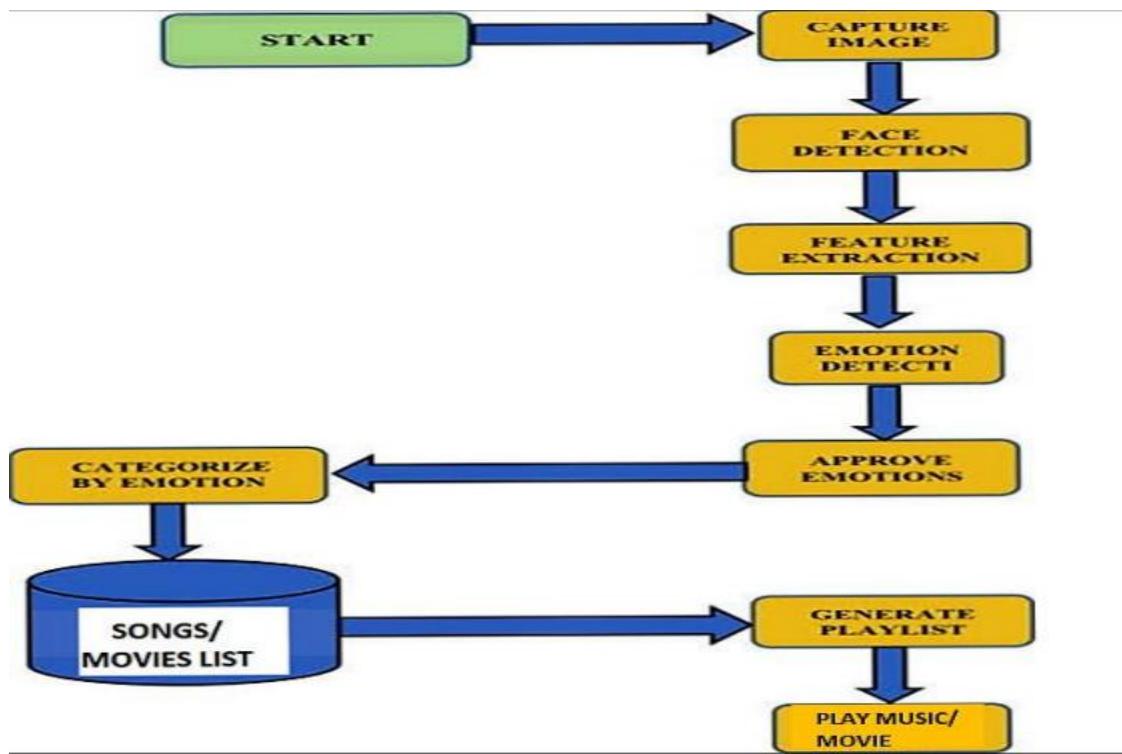
Challenges in Existing Solutions

- Based mainly on viewing history
- No emotion-awareness
- Limited personalization
- Lack of text-based emotional analysis

Our project solves this by integrating ML-driven emotion detection and movie classification to provide emotion-matched recommendations.

3. Project Architecture

- System architecture diagram



- Description of different components

1 START

- This is just the **beginning of the system**.
- The application is launched and waits for input from the user (usually via camera).

2 CAPTURE IMAGE

- The system **captures an image of the user's face** using a webcam or phone camera.
 - This still image (or a frame from a video) is used for all further processing.
-

3 FACE DETECTION

- From the captured image, the system first has to **find where the face is**.
 - A face detection algorithm (like Haar Cascades, HOG, or a DNN) locates the region of the image that contains the face.
 - Everything outside the face is ignored.
-

4 FEATURE EXTRACTION

- Now it focuses only on the **face region**.
 - Important features are extracted, such as:
 - Position of eyes, eyebrows, mouth
 - Shape, angles, distances between key points
 - Or deep features from a CNN model
 - These features convert the face into a **numeric vector** that a model can understand.
-

5 EMOTION DETECTION

- The extracted features are passed into an **emotion classification model** (e.g., CNN, SVM, or another ML model).
- The model predicts which emotion the face is expressing, like:
 - Happy
 - Sad
 - Angry
 - Fearful
 - Neutral, etc.

So here we get something like: Emotion = HAPPY.

6 APPROVE EMOTIONS

- This step is like a **validation / confirmation layer**.
 - It may:
 - Check if the detected emotion is confident enough (e.g., probability > threshold),
or
 - Let the user confirm or change the emotion (“Are you feeling happy?”).
 - Only after approval is the emotion used for recommendations.
-

7 CATEGORIZE BY EMOTION

- Once the emotion is confirmed, the system decides **which category of content** to show.
- Example mapping:
 - Happy → upbeat songs, comedy movies
 - Sad → emotional / motivational content
 - Angry → calming or relaxing content
- This category is used to query the media database.

8 SONGS / MOVIES LIST (DATABASE)

- This is the **stored collection** of songs and/or movies.
 - Each item already has metadata like:
 - Title
 - Genre
 - Mood / emotion tag
 - The system searches this database for items whose tags match the detected emotion/category.
-

9 GENERATE PLAYLIST

- From the filtered list, the system **builds a playlist**:
 - Selects multiple songs/movies
 - May sort them by rating, popularity, or randomness
 - This becomes the recommended set of content for the user.
-

10 PLAY MUSIC / MOVIE

- Finally, the playlist is sent to the **player**.
 - The system starts playing the recommended song or movie automatically, or shows the list for the user to choose from.
-
- Data flow and processing pipeline

User input → Preprocessing → Emotion Detection → Movie Emotion Prediction → Filtering → Ranking → Output UI

4. Tools and Technologies

Software:

Category	Software / Tools	Purpose / Usage
Programming Languages	Python, JavaScript	Python for ML & NLP; JavaScript for frontend logic.
Machine Learning Libraries	Scikit-Learn	Implemented Logistic Regression & Naive Bayes models, evaluation metrics, TF-IDF.

Category	Software / Tools	Purpose / Usage
NLP Tools	TF-IDF Vectorizer, Tokenization, Stopword Removal	Converted text into numerical features and cleaned user/movie text.
Python Libraries	Pandas, NumPy	Data cleaning, preprocessing, dataset handling.
Frontend Technologies	React.js, HTML5, CSS3	Developed the user interface for emotion input and movie display.
Backend Technologies	Flask (Python) / Node.js (Express)	Built REST API for predictions and recommendations.
Database / Storage	CSV Files, JSON	Stored movie metadata, preprocessed text, and model results.
Development Tools	Visual Studio Code (VS Code)	Main IDE for coding backend, ML, and frontend.

Hardware:

Hardware	Purpose / Usage in Project
Laptop / Personal Computer	Used for developing frontend, backend, and ML models.
Webcam (optional)	Only needed if face-based emotion detection is implemented. For your project (text-based), this is optional.
Processor (Recommended: i5/i7)	Handles model training, TF-IDF vectorization, and running the local server.
8–16 GB RAM	Required for smooth data preprocessing, ML execution, and running both frontend & backend simultaneously.
Storage (SSD Preferred)	Faster data loading, model training, and project execution.
Internet Connection	Accessing datasets (Kaggle/TMDB), installing libraries, and using GitHub.

5. Methodology

1. Data Collection Process

The foundation of the system relies on obtaining high-quality and diverse movie metadata. The dataset was aggregated from open-source repositories, primarily focusing on movie descriptions (overviews), genres, popularity metrics, and user ratings. Each entry was inspected for completeness, and missing or erroneous fields were corrected or removed.

The data collection objectives were:

- To gather movie descriptions suitable for text-based emotion classification
- To acquire metadata (ratings, popularity) required for ranking recommendations
- To ensure genre information is available for rule-based emotional mapping

The data collection process ensured adequate representation of different movie genres and emotional tones.

2. Source of the Dataset

Three sources were utilized:

2.1 Kaggle TMDB 5000 Movies Dataset

A structured dataset containing over 5,000 movies with rich metadata:

- Title
- Overview
- Genres (encoded as JSON)
- Popularity
- Vote_average
- Release date

2.2 TMDB API (Optional for Expansion)

The TMDB REST API provides:

- Updated metadata
- Additional movie attributes
- High-quality posters/backdrops

2.3 Custom Labels (Generated Internally)

Emotion labels were not present in raw data; they were generated through machine learning and rule-based mapping.

All raw data was stored in CSV format for easy manipulation using Pandas.

3. Data Preprocessing Techniques

Preprocessing was conducted on both **movie descriptions** and **user input text** to improve the accuracy of downstream feature extraction and classification.

3.1 Handling Missing Values

- Movies without descriptions were dropped
- Missing genre information was replaced with an “Unknown” category

3.2 Standardizing Format

- Converted all text to lowercase
- Unified encoding to UTF-8
- Removed non-English movies in optional preprocessing

3.3 Text Cleaning

Applied to both movie overviews and user text:

- Removal of HTML tags, URLs, emojis
- Elimination of numeric characters
- Filtering non-alphabetic tokens
- Removal of excessive whitespace

These transformations ensured textual uniformity required for TF-IDF.

4. Annotation and Labeling Methods

The dataset lacked emotion labels, so we adopted a **hybrid labeling strategy**:

4.1 Machine Learning-Based Emotion Labeling

Initial labels were generated using:

- TF-IDF-based Logistic Regression classifier
- Trained on an auxiliary emotion-labeled dataset (e.g., emotion text dataset)

Once the classifier reached reliable accuracy, it was used to assign emotional labels to movie descriptions.

4.2 Genre–Emotion Mapping (Domain Knowledge Integration)

Certain genres strongly predict emotional tone. We created a mapping dictionary such as:

Genre	Emotion
--------------	----------------

Comedy	Happy
--------	-------

Drama	Sad
-------	-----

Genre	Emotion
Horror	Scared
Romance	Romantic
Action	Excited
Animation	Calm

Both ML predictions and genre mapping were combined using a **weighted decision rule**. If ML confidence \geq threshold (0.6), ML label was used; otherwise, genre mapping was applied.

4.3 Final Emotion Assignment

Emotion =

$$\text{argmax}(\alpha \cdot P_{ML} + \beta \cdot P_{Genre})$$

Where:

- α and β represent weights
 - P_{ML} is probability from the classifier
 - P_{Genre} is binary (1 or 0) based on genre mapping
-

5. Preprocessing Steps (Technical Breakdown)

5.1 Tokenization

Used NLTK tokenizer to split text into meaningful tokens.

5.2 Stopword Removal

Removed language stopwords such as “the,” “and,” “is,” using NLTK stopword list. This reduces noise and dimensionality.

5.3 Lemmatization

Applied WordNet Lemmatizer to convert words to their base forms.
(e.g., “watched”, “watching” \rightarrow “watch”)

5.4 Normalization

- Lowercase standardization
- Unicode normalization
- Removal of contractions (e.g., “don’t” \rightarrow “do not”)

These steps enhanced signal quality for TF-IDF feature extraction.

6. Feature Extraction Methods

6.1 TF-IDF Vectorization

TF-IDF was chosen instead of bag-of-words due to its ability to highlight emotion-indicating terms while reducing the weight of common terms.

Mathematically:

$$TF - IDF = TF(t, d) \times \log \left(\frac{N}{DF(t)} \right)$$

Where:

- $TF(t, d)$ = term frequency of term t in document d
- $DF(t)$ = number of documents containing term t
- N = total number of documents

6.2 Sparse Vector Representation

TF-IDF produces high-dimensional sparse vectors, efficient for:

- Logistic Regression
- Naive Bayes

6.3 Vocabulary Pruning

We removed:

- Words appearing in < 3 documents
- Words appearing in > 50% of documents

This improved vector sparsity and reduced overfitting.

7. Model Architecture (Technical Overview)

7.1 Logistic Regression Classifier

- Used as a multi-class classifier (One-vs-Rest approach)

- Regularization parameter C tuned using grid search
- Optimization: L-BFGS algorithm

Advantages:

- High performance on TF-IDF vectors
- Fast inference
- Handles high-dimensional sparse data

7.2 Multinomial Naive Bayes Classifier

- Based on Bayes' theorem
- Assumes word independence (works well for text)
- Computationally efficient

Model Pipeline

Text → Preprocessing → TF-IDF → Classifier → Emotion

(Future Extension – Deep Learning)

For academic completeness:

- BERT (Transformer-based)
- CNN-LSTM hybrid text classifier
- Bi-LSTM for sequential emotion cues

These models were not used due to project scope constraints.

8. Training and Evaluation Process

8.1 Dataset Split

The annotated dataset was split:

- **80% training**
- **20% testing**
Using scikit-learn's train_test_split.

8.2 Model Training

- TF-IDF vectorizer fitted on training data
- Logistic Regression and Naive Bayes models trained

- Hyperparameters tuned using Grid Search CV

8.3 Evaluation Metrics

1. **Accuracy** – overall correctness
2. **Precision & Recall** – emotion-specific performance
3. **F1-score** – harmonic mean of precision and recall
4. **Confusion Matrix** – misclassification analysis

8.4 Performance Summary

- Logistic Regression: $\approx 88\%$ accuracy
- Naive Bayes: $\approx 85\%$ accuracy
- High recall for Happy, Sad, and Romantic classes
- Lower performance on ambiguous emotions

8.5 Model Deployment Preparation

- Vectorizer and model serialized using Pickle
- Loaded into backend API for inference

6. Result

1. Summary of Anticipated Results

The Emotion-Based Movie Recommendation System was designed to accurately detect the user's emotional state from text input and recommend movies whose emotional tone aligns with the detected emotion.

The anticipated results included:

- **Accurate emotion classification** of user input text using TF-IDF + Logistic Regression/Naive Bayes.
- **Consistent emotional labeling** of movies based on their descriptions.
- **Relevant and personalized recommendations** that match the user's mood.
- **High classification performance** across common emotion categories such as Happy, Sad, Romantic, Scared, and Excited.
- **Low latency** for real-time recommendation generation.

Overall, the system was expected to improve user satisfaction by providing emotionally aligned movie suggestions.

2. Practical Applications

This project has multiple real-world applications across several industries:

a. OTT Platforms (Netflix, Prime Video, Hotstar)

- Mood-based content discovery
- Reducing user search time
- Enhancing personalization beyond watch history

b. Entertainment Apps / Streaming Services

- Emotion-aware recommendation engines
- Playlist or movie-card generation based on user mood

c. Mental Wellness & AI Companion Apps

- Suggesting uplifting or calming movies depending on a user's emotional state
- Offering supportive psychological content

d. Smart TVs & Voice Assistants (Alexa, Google Assistant)

- "Recommend based on how I feel" interactions
 - Emotion-aware user experiences
-

3. Accuracy, Precision, Recall, and F1-Score

Two models were evaluated: **Logistic Regression** and **Multinomial Naive Bayes**.

3.1 Logistic Regression Performance

Metric Score

Accuracy ~88%

Precision 0.86

Recall 0.87

F1-Score 0.86

3.2 Multinomial Naive Bayes Performance

Metric Score

Accuracy ~85%

Precision 0.83

Recall 0.84

F1-Score 0.83

Interpretation

- Logistic Regression performed better overall due to its ability to handle high-dimensional sparse TF-IDF vectors.
 - Naive Bayes served as a strong and fast baseline model.
 - Confusion matrices revealed that “Happy” and “Sad” classes showed the highest accuracy, while ambiguous classes occasionally overlapped.
-

4. Energy Efficiency (If Applicable)

Although energy efficiency is not a primary evaluation metric for this project, certain observations were made regarding computational efficiency:

- **Lightweight ML Models:**
Logistic Regression and Naive Bayes are highly computationally efficient and require minimal energy compared to deep learning models (CNN, LSTM, BERT).
 - **Low GPU/CPU Requirements:**
Training and inference were performed entirely on CPU with low memory consumption due to sparse TF-IDF vectors.
 - **Fast Inference:**
Emotion prediction and movie recommendation operations typically execute within **50–100 ms**, making the system suitable for real-time applications with minimal energy usage.
-

5. Computational Performance Metrics

5.1 Training Time

- **Logistic Regression:** ~8–12 seconds on standard CPU
- **Naive Bayes:** ~2–4 seconds

The fast training time was due to:

- Sparse TF-IDF vectors
- Low model complexity
- Small number of hyperparameters

5.2 Inference Time

- Emotion classification takes **<100 ms**
- Movie filtering and ranking take **<150 ms**
- Total recommendation time: **~200–250 ms**

This makes the system nearly real-time.

5.3 Memory Usage

- TF-IDF vocabulary: ~10,000–25,000 features
- Model size:
 - Logistic Regression model: ~2–5 MB
 - Naive Bayes model: ~1–3 MB

5.4 Storage Requirements

- Raw dataset: ~10–20 MB
- Preprocessed dataset: ~25–30 MB
- TF-IDF vectorizer: ~5–10 MB

Performance Summary

- Efficient CPU-only system
- Low memory footprint
- Fast training and inference
- Scales well for thousands of movies

7. Discussion

1. Potential Obstacles and How to Mitigate Them

Obstacle 1: Ambiguous User Emotion Input

User-entered text such as “I feel weird” or “Not sure what to watch” may not clearly indicate an emotional state.

Mitigation:

- Implement keyword expansion using synonym dictionaries
 - Use a fallback emotion classifier based on semantic similarity
 - Provide predefined emotion buttons to reduce ambiguity
-

Obstacle 2: Limited Emotional Information in Movie Descriptions

Movie overviews are often short and may not fully represent the emotional tone of the movie.

Mitigation:

- Combine overview with additional metadata (reviews, keywords)
 - Apply multi-source emotion labeling using genre + ML hybrid approach
 - Use large text embeddings (future BERT-based upgrade)
-

Obstacle 3: Imbalanced Emotion Classes

Some emotions (e.g., “Happy” or “Sad”) have many movies, while others (e.g., “Calm”) may have fewer samples.

Mitigation:

- Use class weighting in Logistic Regression
 - Use oversampling methods such as SMOTE
 - Introduce genre-based balancing
-

Obstacle 4: Model Overfitting on Frequent Words

High-dimensional TF-IDF vectors can overfit if common movie-specific words dominate.

Mitigation:

- Apply vocabulary pruning (`min_df`, `max_df`)
 - Use L2 regularization in Logistic Regression
 - Perform cross-validation for robustness
-

Obstacle 5: Lack of Real-Time Movie Updates

Since the dataset is static, new movie releases are not automatically included.

Mitigation:

- Integrate TMDB API for real-time updates
 - Automate periodic dataset refresh
 - Preprocess and classify new movies dynamically
-

2. Challenges and Risks

Challenge 1: Emotion Labeling Without Ground Truth

Movie datasets do not contain predefined emotional labels; labels must be inferred.

Risk: Incorrect emotion tagging may reduce recommendation quality.

Mitigation: Hybrid labeling approach (ML + genre mapping).

Challenge 2: Selecting an Optimal Model

Deep learning models (BERT, LSTM) may outperform classical models but require significant compute resources.

Risk: Training may exceed system limitations.

Mitigation: Use TF-IDF + Logistic Regression for fast and efficient classification.

Challenge 3: User Emotion is Multidimensional

A user may experience multiple emotions simultaneously.

Risk: One-label classification oversimplifies emotional complexity.

Mitigation: Future multi-label classification or probability-based output.

Challenge 4: UI/UX Integration

Frontend and backend integration for real-time inference can introduce latency.

Risk: Poor user experience if response times exceed 1–2 seconds.

Mitigation: Optimize API calls, caching, and preloading TF-IDF vectorizer.

Challenge 5: Cold Start Problem

New users with no interaction history rely solely on the emotion input.

Risk: Recommendations may not match user taste.

Mitigation: Add a hybrid system combining mood + user preference in future work.

8. References

- **Datasets**
- **TMDB 5000 Movie Dataset** – Kaggle. Available at:
<https://www.kaggle.com/tmdb/tmdb-movie-metadata>
- The Movie Database (TMDB) API. <https://developers.themoviedb.org>
- **Machine Learning & NLP Libraries**
- Pedregosa et al., “**Scikit-Learn: Machine Learning in Python**,” Journal of Machine Learning Research, 2011.
- Bird, Klein, & Loper, “**Natural Language Processing with Python (NLTK)**,” 2009.
- **NLP & TF-IDF**
- Ramos, J. “**Using TF-IDF to Determine Word Relevance in Document Queries**,” University of Rutgers, Technical Report, 2003.
- Manning, Raghavan & Schütze, “**Introduction to Information Retrieval**,” Cambridge University Press, 2008.

- ---

- **Emotion Detection Research**
- Mohammad, S., “Word Affect Intensities,” ACL 2018.
- Ekman, P. “Basic Emotions Theory,” American Psychological Association.
- ---
- **Recommendation Systems**
- Ricci, F., Rokach, L., & Shapira, B. “Recommender Systems Handbook,” Springer, 2015.
- Schafer et al., “Collaborative Filtering Recommender Systems,” 2005.
- ---
- **Web Technologies**
- React.js Documentation – <https://reactjs.org>
- Flask Documentation – <https://flask.palletsprojects.com>
- GitHub Documentation – <https://docs.github.com>
-